

# Learning High-Order Word Representations

---

Konstantinos Kogkalidis

June 12, 2018

LoLa Fan Club

# Motivation

---

Idea: structure-preserving map  $\mathcal{F}$

$$\mathcal{F} : \mathcal{G} \rightarrow \mathbf{FdVect}$$

- Atomic types translated to vectors (order-one tensors)
- Complex types translated to (multi-)linear maps (higher order tensors)

## Example

Word Type	$\mathcal{G}$ Type	$\mathcal{F}$ Translation
Noun	$NP$	$\mathbb{R}^{NP}$
Adjective	$NP \backslash NP$	$\mathbb{R}^{NP \times NP} \equiv \mathbb{R}^{NP} \rightarrow \mathbb{R}^{NP}$

$$cat \in \mathbb{R}^{NP}$$

$$black, stray \in \mathbb{R}^{NP \times NP}$$

$$black\ stray\ cat \in \mathbb{R}^{NP}$$

# Why Compositionality?

- Bridging of formal & distributional semantics
- Syntax-informed meaning derivations
- Modeling of functional words
- Formal treatment of ambiguous derivations
- Contextual Disambiguation
- Richer representations
- $\vdots$

# Why Not Compositionality?

- ✓ Great properties
- ? How to obtain word representations?

Possible options:

# Why Not Compositionality?

- ✓ Great properties
- ? How to obtain word representations?

Possible options:

1. Co-occurrence statistics

# Why Not Compositionality?

✓ Great properties

? How to obtain word representations?

Possible options:

1. Co-occurrence statistics ✗
2. Unsupervised techniques (*a la word2vec*)



# Why Not Compositionality?

✓ Great properties

? How to obtain word representations?

Possible options:

1. Co-occurrence statistics ✗
2. Unsupervised techniques (*a la word2vec*) ✗
3. Supervised learning ?

# Problem Statement

Examine whether supervised learning can be used to find higher-order word representations (transitive verbs)

# Supervised Learning

---

- Search over set of functions  $A \rightarrow B$  parameterized over  $P$
- Find optimal approximation  $\hat{f}_P$  to  $f: A \rightarrow B$
- Use samples  $(a, f(a)) \in A \times B$  to update  $P$

# **Supervised Learning**

---

**Dataset**

Sample space must be:

- labeled
- constrained
- of large size
- of high quality

Raw text paraphrase pairs

## Example pair

*proposed by the president ~ suggested by the chairman*

- labeled ✓
- constrained ✗ (different syntactic types)
- of large size ✓
- of high quality ?

## 1. **Parse and filter by type** (transitive verb case)

- labeled ✓
- constrained ✓
- of large size ✗ (>95% loss)
- of high quality ✗ (*parser-induced errors*)



## 2. **Back-translation**

- Labeled ✓
- constrained ✓
- of large size ✓
- of high quality ✗ (*translation-induced errors*)

### 3. **Filter by co-occurrence / mutual information**

- Labeled ✓
- constrained ✓
- of large size ✓
- of high quality ?

Verb / object dictionaries:

$$\mathcal{V} : \{v_1 : 1, v_2 : 2, \dots, v_N : N\}$$

$$\mathcal{O} : \{o_1 : 1, o_2 : 2, \dots, o_M : M\}$$

Paraphrase relation:

$$\mathcal{P} : \mathbb{N} \times \mathbb{N} \times \mathbb{N} \times \mathbb{N} \rightarrow \{0, 1\} \quad (\text{binary classification})$$

$$\mathcal{P}(i, j, k, l) = \mathcal{P}(k, l, i, j) = \begin{cases} 1 & v_i o_j \sim v_k o_l \\ 0 & \text{otherwise} \end{cases}$$

# **Supervised Learning**

---

## **Formulating the Network**

# Training Objective

Our semantic interpretations are:

- Actions:  $[a] = \mathbb{R}^A$
- Objects:  $[np] = \mathbb{R}^{NP}$
- Transitive Verbs:  $[a/np] = \mathbb{R}^{A \times NP}$

And our objective is to learn a **verb embedding function**  $\epsilon_{verb}$ :

$$\epsilon_{verb} : \mathbb{N} \rightarrow \mathbb{R}^{A \times NP}$$

But instead we have samples from some  $f : \mathbb{N}^4 \rightarrow \{0, 1\}$

## Solution

Formulate  $f_p$  to incorporate  $\varepsilon_{verb}$ .

$$f_p = f_1 \circ f_2 \circ \dots \circ \varepsilon_{verb} \circ \dots$$

## Solution

Formulate  $f_p$  to incorporate  $\varepsilon_{verb}$ .

$$f_p = f_1 \circ f_2 \circ \dots \circ \varepsilon_{verb} \circ \dots$$

## Simplification (1)

Assume pre-trained **object embedding function**  $\varepsilon_{object}$ .

$$\varepsilon_{object} : \mathbb{N} \rightarrow \mathbb{R}^{300}$$

# Filling the missing blocks

$$i \in \mathbb{N}$$

$$j \in \mathbb{N}$$

$$k \in \mathbb{N}$$

$$l \in \mathbb{N}$$

$$\hat{y} \in \mathbb{R}$$



# Filling the missing blocks

$$i \in \mathbb{N}$$

$$j \in \mathbb{N}$$

$$k \in \mathbb{N}$$

$$l \in \mathbb{N}$$

$\varepsilon_{object}$


$$j \in \mathbb{R}^{300}$$

$\varepsilon_{object}$

$$l \in \mathbb{R}^{300}$$


$$\hat{y} \in \mathbb{R}$$

# Filling the missing blocks



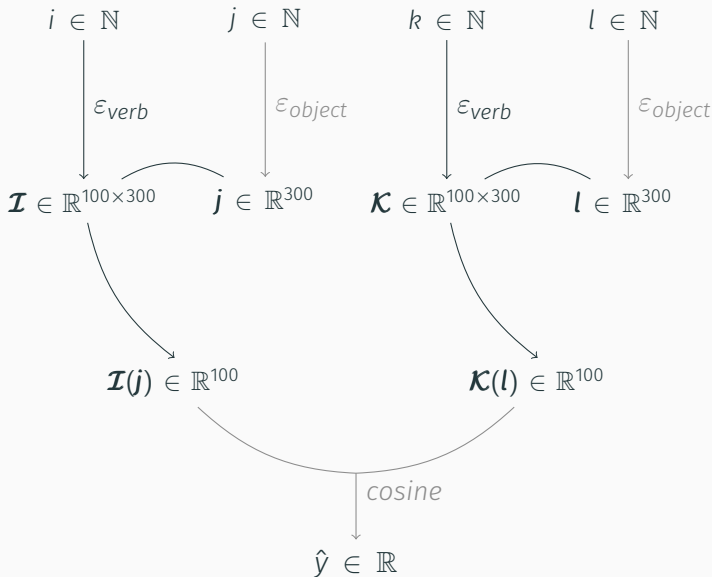
$$\hat{\mathbf{y}} \in \mathbb{R}$$

# Filling the missing blocks



$$\hat{y} \in \mathbb{R}$$

# Filling the missing blocks



## Objective Function

$$\cos(\mathbf{V}_i(j), \mathbf{V}_k(l)) \rightsquigarrow \mathcal{P}(i, j, k, l) \quad \forall (i, j, k, l) \in \mathcal{V} \times \mathcal{O} \times \mathcal{V} \times \mathcal{O}$$

## Objective Function

$$\cos(\mathbf{V}_i(j), \mathbf{V}_k(l)) \rightsquigarrow \mathcal{P}(i, j, k, l) \quad \forall (i, j, k, l) \in \mathcal{V} \times \mathcal{O} \times \mathcal{V} \times \mathcal{O}$$

Considerations:

### 1. Network Size

- 1.000 verbs
  - $100 \times 300 = 30.000$  parameters per verb
- $\Rightarrow$  30.000.000 parameters for  $\varepsilon_{verb}$  to learn

## Objective Function

$$\cos(\mathbf{V}_i(j), \mathbf{V}_k(l)) \rightsquigarrow \mathcal{P}(i, j, k, l) \quad \forall (i, j, k, l) \in \mathcal{V} \times \mathcal{O} \times \mathcal{V} \times \mathcal{O}$$

Considerations:

### 1. Network Size

- 1.000 verbs
  - $100 \times 300 = 30.000$  parameters per verb
- $\Rightarrow$  30.000.000 parameters for  $\varepsilon_{verb}$  to learn

### 2. Quantifying over two spaces ...

## Objective Function

$$\cos(\mathbf{V}_i(j), \mathbf{V}_k(l)) \rightsquigarrow \mathcal{P}(i, j, k, l) \quad \forall (i, j, k, l) \in \mathcal{V} \times \mathcal{O} \times \mathcal{V} \times \mathcal{O}$$

Considerations:

1. Network Size

- 1.000 verbs
  - $100 \times 300 = 30.000$  parameters per verb
- $\Rightarrow$  30.000.000 parameters for  $\varepsilon_{verb}$  to learn

2. Quantifying over two spaces ...

3. ... both of which are non-convex





... A "beast" to train ☹️

# **Supervised Learning**

---

## **Transferring Knowledge**

## Finding an Oracle

Dataception: use our labeled dataset to create a new dataset

# Finding an Oracle

Dataception: use our labeled dataset to create a new dataset

## Simplification (2)

Assume another pre-trained verb embedding function  $\epsilon'_{verb}$ .

$$\epsilon'_{verb} : \mathbb{N} \rightarrow \mathbb{R}^{300}$$

# Finding an Oracle

Dataception: use our labeled dataset to create a new dataset

## Simplification (2)

Assume another pre-trained **verb embedding function**  $\epsilon'_{verb}$ .

$$\epsilon'_{verb} : \mathbb{N} \rightarrow \mathbb{R}^{300}$$

## Oracle

We can now train a **paraphrase embedding function**  $\epsilon_{par}$ .

$$\epsilon_{par} : \mathbb{R}^{300} \times \mathbb{R}^{300} \rightarrow \mathbb{R}^{100}$$

# Oracle Flow

$$i \in \mathbb{N}$$

$$j \in \mathbb{N}$$

$$k \in \mathbb{N}$$

$$l \in \mathbb{N}$$

$$\hat{y} \in \mathbb{R}$$

# Oracle Flow

$$i \in \mathbb{N}$$

$$j \in \mathbb{N}$$

$$k \in \mathbb{N}$$

$$l \in \mathbb{N}$$

$\varepsilon_{object}$


$$j \in \mathbb{R}^{300}$$

$\varepsilon_{object}$

$$l \in \mathbb{R}^{300}$$

$$\hat{y} \in \mathbb{R}$$


# Oracle Flow



$$\hat{y} \in \mathbb{R}$$

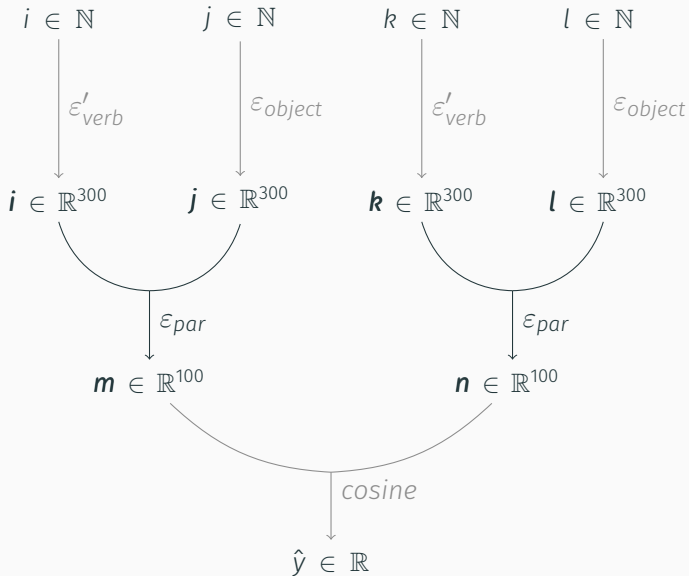


# Oracle Flow




$$\hat{y} \in \mathbb{R}$$

# Oracle Flow



# Training the Oracle

$\epsilon_{par}$ : recurrent autoencoder ( $\approx 700.000$  parameters)




## New Objective Function

$$\cos(\mathbf{V}_i(j), \varepsilon_{par}(\mathbf{v}, j)) \rightsquigarrow 1 \quad \forall (i, j) \in \mathcal{V} \times \mathcal{O}$$


- $\varepsilon_{par}$  gives us paraphrase embeddings '*for free*'
- We can use them to facilitate training
- Much smaller problem space

# Composing Networks



# Training the Original

$\epsilon_{verb}$ : tanh activated dense layer ( $\approx 30.000.000$  parameters)





**The beast has been tamed! 😊**

# Evaluation


---



Task-specific performance relates to the small-scale structure of the learned space:

Ground Truth \ Prediction	Oracle		Final	
	T	$\perp$	T	$\perp$
T	0.92	0.08	0.88	0.02
$\perp$	0.08	0.92	0.12	0.98

## 3D PCA on paraphrase embeddings



# Conclusion

---

1. Uninformative error signal ?
2. Over-parameterization
  - a) Encoder / decoder architectures ?
  - b) Linearity constraint ✗
  - c) Chasing after an oracle ?
  - d) Bad scaling ✗

## Next Steps

1. Directly evaluate verb matrices
2. More structural constraints (activity regularization)
3. Iterative learning
4. Other data formats:
  - Different syntactic types
  - Different labels / samples altogether
5. Different oracle architectures

# Bag of Tricks

---

1. Negative Sampling
2. Curriculum Learning / Domain Adaptation
3. Loss Mixing