

Transfer Learning High-Order Word Representations

Konstantinos Kogkalidis

konstantinos@riseup.net

1 Preface

This document is meant to act as a brief technical overview of my work on transfer-learning high-order tensorial representations for transitive verbs. As the project itself is still work in progress, the document will also be evolving over time, tracking issues and solutions as they appear.

1.1 Changelog

Date	Changes
23/06/2018	First Draft Preface Introduction Intractability

2 Introduction

Distributional Compositional Semantics are an active area of computational linguistics research, concerned with the modeling of word interactions to form larger linguistic units. Inspired by formal semantics and distributional models of meaning, it proposes a functional merging of the two. The core idea is that, given a grammar formalism \mathcal{G} , such as pregroup grammars or combinatory categorial grammars, we may move to the category of finite-dimensional vector spaces, \mathbf{FdVect} via a structure-preserving map (homomorphism) $\mathcal{F} : \mathcal{G} \rightarrow \mathbf{FdVect}$.

The grammar is responsible for deriving proofs for well-formed linguistic units. In such a grammar, words are assigned types descriptive of their syntactic role. These types may either be atomic (e.g. s for sentences or NP for noun-phrases) or complex, defined as the structured compositions of atomic ones (e.g. NP/NP for adjectives, "consuming" a noun-phrase to their right to produce a new, larger noun-phrase).

Using \mathcal{F} , we may assign vector spaces to simple types of \mathcal{G} . Complex types can then be translated to high-order tensors (or equivalently multi-linear maps), which operate on their arguments (as dictated by \mathcal{G}), to produce syntactically-informed vectorial constructions for larger phrases. Crucially, this allows us to benefit from both the applied neuro-distributional approach and the more theoretically grounded aspects of computational linguistics in meaningfully constructing phrase or sentence numerical representations.

Word Type	\mathcal{G} Type	\mathcal{F} Translation
Noun	NP	$[NP]$
Adjective	NP / NP	$[NP] \rightarrow [NP]$

Table 1: Example translations from grammatic types to \mathcal{G} types to vector spaces. If noun-phrases occupy an arbitrary vector space, $[NP]$, then adjectives are an endomorphism of that space.

A key problem in the field is the lack of an existing, experimentally verified, framework towards obtaining these high-order tensors for complex word types. This work proposes an alternative to thus far explored solutions, namely transforming the problem setting to that of supervised learning.

3 Towards a Supervised Setting

3.1 Overview

Supervised Learning can be abstractly summarized as the search process over a set of continuously differentiable functions F_p with predetermined signature $A \rightarrow B$, parameterized over a set of trainable weights p

$$F_p = \{f_p : A \rightarrow B\}$$

with the goal being finding the one function \hat{f}_p that best approximates some ideal (but unknown) function $g : A \rightarrow B$.

3.2 Dataset

To begin a supervised learning endeavour, we need to first find samples from the aforementioned ideal function g , organized in a tuple $(X, g(X))$. Finding such samples and evaluating their potential to drive learning is by itself a difficult task. For this work, we will be working with the Paraphrase Database (PPDB) phrasal corpus. This contains a large amount of phrases organized in paraphrastic pairs, i.e. phrases that are semantically close to each other.

Before actually utilizing the dataset, a lot of preprocessing was necessary for a variety of reasons summarized below.

Parsing Working with any potential syntactic type would require a dynamically adaptive graph, a functionality that is only being provided by a small number of relevant libraries on an experimental level. Additionally, it would incur heavy computational costs (as explained in Section 5) and would make debugging and evaluating performance hard. For these reasons, it was deemed a minor compromise to constrain the experimentation domain on just transitive verbs and their objects, and the question we set out to answer is whether we can find a way to learn meaningful verb matrices (i.e. order-2 tensors). We use a (very) simplified grammar consisting of two simple types:

1. **Noun Phrases** NP

2. **Actions** A

Where the type A encompasses actions potentially undertaken by some subject agent. Then a mono-transitive verb, which is the focal point of this work, can be seen as a the complex type $[a/np]$ and its \mathcal{F} translation as a morphism from the space of noun phrases to that of actions. Under this light, we define and use the below type translations:

1. **Objects** $[np] = \mathbb{R}^{300}$

2. **Actions** $[a] = \mathbb{R}^{100}$

3. **Monotransitive Verbs** $[a/o] = \mathbb{R}^{60 \times 300}$

To enable experimentation, all of our samples need to follow this single derivation. This is accomplished by filtering the entirety of the PPDB corpus by spaCy’s syntactic parser. The choice of the concrete vector spaces used is somewhat arbitrary but justifiable. Objects are chosen to live in \mathbb{R}^{300} since that is the standard word vector dimensionality used by popular pretrained word embedding systems such as word2vec, GloVe and fastText, thus giving us the ability to inherit these vectors without further training. The action space is set to \mathbb{R}^{100} ; fewer dimensions would be overly compressive, reducing the capacity for disambiguation and prediction, whereas higher dimensions would make the network more prone to overfitting and harder to train (more on Section 5).

Backtranslation The syntactic filtering vastly reduces the number of available samples. A large number of different verbs, but also different verb-object combinations, is necessary for training to succeed. Consequently, we need a way to recover from this sample loss. A good way around the problem is to use backtranslation, a popular technique for generating synthetic new paraphrastic phrases out of existing ones. The procedure involves using a pretrained machine translation black box which allows translating an arbitrary phrase to and from a set of languages. For the purposes of this work, google cloud services' translation was used. Given such a system, we can initiate a (potentially randomized) set of candidate languages and draw a randomized transition graph (potentially including loops) passing through those. The stochastic and slightly inaccurate nature of neural translation and the expressive ambiguity between different languages are guaranteed to introduce a minor semantic shift with each translation. After each phrase is chain-translated up to the graph's end-point, we can reverse translate back to the original language in a single step. This usually results in a phrase that is paraphrastic to our starting one, yet not identical to it. This process can be repeated until our dataset size is sufficient for training.

Statistical Filtering PPDB is itself automatically annotated, meaning that it contains a significant amount of error. Syntactic parsing and backtranslation are also not 100% accurate, therefore introducing additional errors with each step of the process. This error is hard to control; manually checking the entirety of the samples would be unfeasibly demanding. A cost-efficient way of reducing the error to acceptable rates is to apply statistical filtering on our samples. Since we are seeking to learn verb representations, it would make sense to only keep verbs that appear adequately often in our samples, yet are not extremely frequent (thereby carrying more noise than information). The distribution of unique verbs in our dataset follows a very sharp, uneven distribution; the head and tail of these are truncated to even out verb frequency. Additionally, verbs co-existing with a very small number of objects carry little information and are prone to overfitting against their objects and are also cut out.

Machine Readability The last preprocessing step is to make the dataset machine-readable; we initiate two injective dictionaries, mapping verbs and objects to and from unique naturals:

$$\begin{aligned}\mathcal{V} &: \{v_1 : 1, v_2 : 2, \dots, v_N : N\} \\ \mathcal{O} &: \{o_1 : 1, o_2 : 2, \dots, o_M : M\}\end{aligned}$$

Via this encoding, we can now formulate our ideal function as a relation \mathcal{P} ; it is a function that accepts two tuples of two naturals each (one tuple uniquely determines a verb-object pair), and returns either 1 if the two encoded phrases are paraphrastic or 0 otherwise:

$$\begin{aligned}\mathcal{P} &: \mathbb{N} \times \mathbb{N} \times \mathbb{N} \times \mathbb{N} \rightarrow \{0, 1\} \\ \mathcal{P}(i, j, k, l) &= \mathcal{P}(k, l, i, j) = \begin{cases} 1 & v_i o_j \sim v_k o_l \\ 0 & v_i o_j \not\sim v_k o_l \end{cases}\end{aligned}$$

The verb(/object) identifiers can be one-hot encoded before being fed as an input to any embedding layer. One-hot encoding translates a natural n from a strictly ordered set of size N to a N -dimensional vector, whose only non-zero element is its n -th one (being equal to 1). This allows each verb(/object) to be assigned a tensorial representation independent of how the rest of the words are represented. The output is already in an ideal form for a classification task.

4 Student Network

Given the ideal function \mathcal{P} , we may begin assembling a network. The goal we set was to train some verb embedding function, which takes a verb's identifier and returns its correspondig matrix:

$$\varepsilon_{verb} : \mathbb{N} \rightarrow \mathbb{R}^{100 \times 300}$$

Including the one-hot encoding process, the function is simply a trainable linear map implemented by a feedforward, linearly activated, single-layer network with no bias and weights:

$$\mathbf{E}_{verb} \in \mathbb{R}^{|V| \times 30,000,000}$$

followed by a reshape. For reasons of brevity, the one-hot encoding and reshaping processes will be left implicit from now on.

The function we wish to train, ε_{verb} , seems to have little in common with our ideal function's \mathcal{P} signature.

5 Objective Function

The network, as depicted above, type checks and is fully differentiable. Given some binary loss operator $\mathcal{L}[\cdot, \cdot]$, the objective function can be formulated as:

$$\min_{\forall i,j,k,l} \mathcal{L}\{\cos(\mathbf{V}_i(\mathbf{o}_j), \cos(\mathbf{V}_k, \mathbf{o}_l)), \mathbf{P}(i, j, k, l)\}$$

Potential candidates for \mathcal{L} might be mean square error, binary crossentropy, categorical hinge or any linear combination thereof. A short discussion on how each particular choice affects the training process can be found in Section todo.

Disregarding the particular choice of \mathcal{L} , let us translate the objective function into natural language. We want to minimize some measure of classification penalty between the paraphrase prediction of two phrases and their original label, as given by our training set. More concretely, we want to obtain a phrase embedding from each phrase, constructed by applying an object vector to a verb matrix, such that when we compare two phrase vectors against one another via cosine similarity the result is close to 1 if these two phrases are labeled as paraphrastic or 0 otherwise.

5.1 Intractability

The complexity of the raw optimization problem should be apparent by its formulation alone. The objective is in fact intractable, i.e. practically impossible to optimize over. To better understand why, we need to address a number of key points.

Number of parameters Our verb embedding function is a fully-connected feedforward layer. Fully connected layers have a trainable multiplicative weight associating each input neuron with each output neuron. In our case, the input is the one-hot encoding of each verb (i.e. a $|V|$ -dimensional vector space) and the output space is that of linear functions $\mathbb{R}^{NP} \rightarrow \mathbb{R}^A$ (i.e. $(NP \times A)$ -dimensional representations). Given the fact that the dataset contains $\sim 1,000$ unique verbs, and our choice of spaces for the objects and actions, the resulting layer enumerates 30,000,000 trainable parameters.

Quantifying Over Multiple Spaces The objective function quantifies over two directly interacting spaces; the space of verbs and the space of actions. This is problematic in many ways; essentially, each space will be sensitive to perturbations of the other (at best), or their joint optimization might be infeasible due to one's minimum being at odds to the other's (at worst).

Non-Convexity More importantly, however, both of the two spaces that we are optimizing over are non-convex. This means that each of them may have a non-unique global minimum and possibly a large number of local minima. Intuitively, the paraphrase space requires that the vectors of two non-paraphrastic phrases are perpendicular to one another. In a two dimensional space, a simple rotation operation applied on all the vectors would still yield a new solution from an existing one. This is even further attenuated in a higher-dimensional space. At the same time, the verb space requires that an object

vector \mathbf{o} applied on a verb matrix \mathbf{V} would result in some predefined phrase vector \mathbf{p} . Analytically solving $\mathbf{V}\mathbf{o} = \mathbf{p}$ w.r.t. to \mathbf{V} is impossible, as there exist infinite possible linear combinations of \mathbf{o} that result in \mathbf{p} . More realistically, we have a system of linear equations $\mathbf{V}\mathbf{o}_1 = \mathbf{p}_1, \dots, \mathbf{V}\mathbf{o}_k = \mathbf{p}_k$ for all k unique objects vectors applied on the same verb \mathbf{V} yielding different phrase vectors, thus constraining our solution space; however, there is still limited guarantee of the solutions being unique.

The above are all clear indicators of the intricacies of attempting to train the network. Even though our operations are end-to-end differentiable, theoretically allowing gradient-based methods to work, the solution space is so vast, interdependent and noisy that we have no guarantee of convergence.